

# Optimization and Profiling

Computing Methods for Experimental Physics and Data Analysis

A. Manfreda

alberto.manfreda@pi.infn.it

INFN-Pisa

Compiled on October 31, 2019



# Optimization and profiling

- ▷ Optimization is the process of making your code more performant
- ▷ Profiling is the study of the resource cost of your code (CPU, RAM, Hard Disk, network . . . ) possibly as function of time / input / execution status.
- ▷ Most important message from today's lesson:
  - ▷ **Premature optimization is the root of all evil**
- ▷ The correct workflow:
  - ▷ Write correct, redible, debugged, tested, documented code
  - ▷ Ignore performance at all
  - ▷ If the performance are ok: great, you are done!
  - ▷ If the performance are not ok: **profile** to find the bottlenecks.
  - ▷ Start the optimization from the problematic lines! Everything else is a waste of time.
  - ▷ Keep testing whenever you have a new version!



## Basic of PC architecture

- ▷ **Disclaimer:** This is a naive description. More on this in the future lessons!
- ▷ Data are processed by CPUs (or GPU)
- ▷ CPUs have a small memory *cache* (or more of them: L1, L2, ...) of a few KB from which it gets the data
- ▷ Generally L1 is on the processor, L2 or L3 are shared among different processors, but this scheme may vary
- ▷ The connection between CPU and cache is called *bus backside* and is very fast
- ▷ Your data needs to be transferred from the RAM to the cache, which happens through the *frontside bus*, slower than the backend bus
- ▷ Sometimes you need to retrieve data from the hard disk or the network: in that case the time required to retrieve the data is even (much) larger



## General optimization goals

- ▷ As obvious as it is: do less operations! Use an optimal algorithm for doing your job
- ▷ Reduce as much as possible the slow operations involving hard drives and networks
- ▷ Make sure that the next data required by the CPU are already in the cache, and need not to be retrieved from the RAM through the slower frontside bus
- ▷ If possible use *vectorization*, that is let the CPU do multiple operations at once
- ▷ If the system has more than one processor, as most computers nowadays, try to parallelize the work on more of them - this will be covered in the next lessons!
- ▷ Python is a high level language and generally does not allow a direct control over the memory usage, or the cache - however there are ways to improve the speed of the code, as we will see!



## What to avoid

- ▷ The CPU uses algorithms of *branch prediction* and *pipelining* in order to try to load the next instructions (and the required data) while processing the current one
- ▷ Whenever that fails you will get *branch misses* and/or *cache misses*, plus usually a number of *stalled cycle* on the frontside or backside bus
- ▷ Using data structures that keep data contiguously in memory is better, as sparse data are more difficult to move into the cache into a single transfer
- ▷ That's why a list is worse than an array or a numpy array for numerical operations
- ▷ *context switches* and *cpu migrations* are managed by the OS, so there is not much you can do about that
- ▷ Memory allocations are also expensive, as the program is paused and wait for the OS to find a free memory location (this, together with I/O operations, is a typical case where a context switch may happen)



## Tools that we will use today

- ▷ Time profiling:
  - ▷ Simple 'print' statements
  - ▷ A time measuring decorator
  - ▷ The *timeit* Python module
  - ▷ The Unix *time* utility
  - ▷ *cProfile*
  - ▷ *line\_profiler*
- ▷ Bytecode study:
  - ▷ *dis*
- ▷ CPU efficiency:
  - ▷ *perf*
- ▷ Memory profiling:
  - ▷ *heapy*



# Introducing the Julia set

Example taken from the book *High Performance Python*

- ▷ Fractal set named after the mathematician Gaston Julia
- ▷ Take a complex function  $f(z)$  and a real number  $R$
- ▷ Apply repeatedly  $f$  to a complex number  $z$ . If the norm of the result is always smaller than  $R$  the number belongs to the set
- ▷ In practice we can only test up to a number of iterations: after that the number will be considered as belonging to the set
- ▷ The function we will use is  $f(z) = z^2 + c$ , where  $c$  is the constant complex number  $c = -0.62772 - 0.42193i$

